

Tutorial 1

1. What measure do we use when analysing performance of out-of-core algorithms? Why? What idealising assumption do we make by doing so?
2. In algorithm analysis, the performance of the algorithm is usually measured as a function of the size of input. For out-of-core algorithms, we measure the performance as a function of the size of the input, N , and the size of the available RAM, B . In what units are N and B measured? Why?
3. What should the relationship between the size of the RAM and the size of the input file be for the sorting to be done in at most 2 passes?
4. Assume that the page size is equal to the number of bytes able to hold two integers; further assume that the RAM size is equal to the number of bytes able to hold 8 integers. Given these assumptions, step through the external merge-sort algorithm applied to the input files containing the following sequences of numbers and determine how many passes are needed to sort these files:
 - (i) 10, 7, 1, 8, 27, 8, 7, 9, 3, 53, 45, 6, 2, 100, 52, 12, 100, 101, 21, 15, 7, 23, 1;
 - (ii) 17, 15, 3, 2, 3, 2, 17, 19, 19, 1, 7, 21, 5, 16, 15, 11, 7, 22, 19, 3.
5. Consider the following 3 scenarios: you have
 - (a) a file whose size equals 10,000 pages and RAM whose size equals 3 pages;
 - (b) a file whose size equals 20,000 pages and RAM whose size equals 5 pages;

- (c) a file whose size equals 2,000,000 pages and RAM whose size equals 17 pages.

For each of the scenarios above, answer the following questions about the external merge-sort algorithm (do not worry about the exact number, perform as much of the calculation as you can without the use of a calculator):

- (i) How many runs will the algorithm produce in the initial pass?
 - (ii) How many passes will it take to sort the entire file?
 - (iii) What is the total cost of sorting the file?
 - (iv) What is the size of the memory needed to sort the entire file in two passes?
6. Consider a relation containing information about university students with the following schema:

```
students (sid: integer, sname: varchar(50),  
street: varchar(50), city: varchar(30), age: integer)
```

Assume that the students file consists of 5,000 pages.

- (i) Let B be the number of pages that fits into the available RAM. If $B = 50$, how many passes (including pass 0) of the external sort-merge algorithm will be required to sort the relation and how many runs will be produced by each of the passes?
 - (ii) For the set-up from the previous question, how many I/Os will be required? Assume that the relation is originally on disk and that the result of the sorting operation must also be written to disk.
 - (iii) Now, consider the query `SELECT sid, age FROM Students ORDER BY age`. Briefly describe how we could execute this query in a way that would result in the lowest possible cost of sorting the `Students` relation? Does your approach save disk reads, disk writes, or both?
7. Consider a simple (non-recursive), 2-pass disk-based hashing strategy described in the lecture, where in the first pass the file is partitioned and in the second pass each of the partitions is hashed.

- (i) Ignoring any space overhead for building hash tables and assuming a perfect hash function, what is the minimum size of RAM needed to hash the above **students** relation in 2 passes?
- (ii) For the previous question, how many I/Os will be required? Assume that the relation is originally on disk and that the result of the hashing operation must also be written to disk.